



AUTOMOTIVE  
LINUX SUMMIT  
EUROPE

# Debugging Tools and Techniques for Virtualized Automotive Systems

**Hiroyuki Ishii**, *Panasonic Automotive Systems Co., Ltd.*  
*AGL System Architect Team & Virtualization Expert*

#EMBEDDEDOSSUMMIT

# Agenda

- Introduction
- Debugging Tools for Virtualized Systems
- Example: Analyzing vhost-net
- Summary

# Agenda

- Introduction
- Debugging Tools for Virtualized Systems
- Example: Analyzing vhost-net
- Summary

# Who Am I

- Hiroyuki Ishii
- > 15 years of experience in software development for IVI and CDC products at Panasonic Automotive Systems
- Expertise: Linux kernel, performance engineering, virtualization, cloud-native technologies
- Joined AGL project since 2021 as:  
Steering Committee, System Architect, Virtualization Expert etc.



# Background Trends in Automotive Software

## Various requirements arising related to Software-Defined Vehicles

- High-performance computing
- Hardware scalability
- Workload distribution
- Mixed criticality

## Growing demand and initiatives for virtualization as a key technology

- AGL actively worked on Virtual Machine (VM) based demo integration
- SOAFEE, Eclipse SDV and Android Automotive are also actively contributing

# Challenges: Complexity of Virtualized Systems

- Complex interactions between host and guest components
- Increased system footprint, complicated integration
- Limited capabilities/visibility within guest environments



- Debugging/performance engineering become increasingly difficult
- Complications in understanding system behavior



Needs for specialized tools and techniques for virtualized systems

# Agenda

- Introduction
- Debugging Tools for Virtualized Systems
- Practice: Analyzing vhost-net
- Summary

# Tools Available for Virtualized Systems

- `perf`
  - `perf stat, perf kvm stat|record`
- `trace-cmd`
  - `trace-cmd record -e kvm`
- `bcc`
  - `virtiostat, kvmexit, ...`
- `flamegraph.pl`
- `debuginfod`



# perf

- A powerful profiling/tracing tool for Linux system
- Frontend for perf\_events, ftrace, kprobes etc.
- Developed within the Linux kernel source tree
  - Stable and reliable
  - Has strong dependency on kernel version
- Some hardware events are platform dependent

# trace-cmd

- Frontend tool for ftrace
  - Developed by Steven Rostedt
- Provides features similar to perf
  - Simpler and more user-friendly (IMO)
- Recent updates designed for embedded/virtualized use cases
  - trace-cmd agent
  - trace-cmd listen
- <https://github.com/rostedt/trace-cmd>

# bcc


- A BPF-based tracing tool for performance analysis and debugging
- Offers rich set of features:
  - Python support
  - Various practical example scripts included
  - Simplifies creation of custom tools built upon it
- Drawbacks:
  - Relatively large dependency (LLVM, Python)
  - Limited support for the ARM platform
- <https://github.com/iovisor/bcc>

# flamegraph.pl

- A visualization tool for performance data samples
- Helps understanding workload overview and identifying bottlenecks
- <https://github.com/brendangregg/FlameGraph>

```
qemu-system-aar 2822 [005] 105761.404072: 8063588 cycles:
ffff8000102cdb44 get_page_from_freelist+0x174 ([kernel.kallsyms])
ffff8000102cfa2c __alloc_pages_nodemask+0x16c ([kernel.kallsyms])
ffff8000102ed824 alloc_pages_current+0x94 ([kernel.kallsyms])
ffff8000102c9970 __get_free_pages+0x20 ([kernel.kallsyms])
ffff800010348020 __pollwait+0x60 ([kernel.kallsyms])
ffff8000103981d8 eventfd_poll+0x68 ([kernel.kallsyms])
ffff800010349868 do_sys_poll+0x268 ([kernel.kallsyms])
ffff800010349e5c __arm64_sys_ppoll+0xac ([kernel.kallsyms])
ffff800010028fa4 el0_svc_common.constprop.0+0x84 ([kernel.kallsyms])
ffff800010029104 do_el0_svc+0x34 ([kernel.kallsyms])
ffff8000111cdb40 el0_svc+0x20 ([kernel.kallsyms])
ffff8000111ce0f4 el0_sync_handler+0xa4 ([kernel.kallsyms])
ffff800010012700 el0_sync+0x180 ([kernel.kallsyms])
ffff800010012700 ppoll+0x98 (/usr/lib/libc.so.6)
aaaad028fd9c main_loop_wait+0x14c (/usr/bin/qemu-system-aarch64)
aaaad00151a8 qemu_main_loop+0x144 (/usr/bin/qemu-system-aarch64)
aaaacfad34 main+0x14 (/usr/bin/qemu-system-aarch64)
ffffa2d4b230 __libc_start_call_main+0x60 (/usr/lib/libc.so.6)
ffffa2d4b30c __libc_start_main@@GLIBC_2.34+0x9c (/usr/lib/libc.so.6)
aaaacfae4cf0 _start+0x30 (/usr/bin/qemu-system-aarch64)

qemu-system-aar 2829 [007] 105761.407769: 2762147 cycles:
ffff80001025e380 clear_rseq_cs+0x20 ([kernel.kallsyms])
ffff80001001e74c do_notify_resume+0x14c ([kernel.kallsyms])
ffff8000100131dc work_pending+0xc ([kernel.kallsyms])
ffffa2dffe90 __GI___ioctl+0x10 (/usr/lib/libc.so.6)
aaaad00ff06c kvm_cpu_exec+0xac (/usr/bin/qemu-system-aarch64)
aaaad0100630 kvm_vcpu_thread_fn+0xb0 (/usr/bin/qemu-system-aarch64)
aaaad0272404 qemu_thread_start+0xa0 (/usr/bin/qemu-system-aarch64)
ffffa2da02a8 start_thread+0x2c8 (/usr/lib/libc.so.6)
ffffa2e07c1c thread_start+0xc (/usr/lib/libc.so.6)
```



(Continues over 100k lines)

# debuginfod

- A host daemon for handling/managing debuginfo files
  - Part of the elfutils project ([git://sourceware.org/git/elfutils.git](https://sourceware.org/git/elfutils.git))
  - Compatible with numerous popular debugging tools (gdb, perf, trace-cmd, bcc, and many others)
- Beneficial for embedded systems
  - Allows for offloading debuginfo files from the target system
  - Resolves path conflicts of debuginfo within virtualized environments
- Integrated in Yocto build system (oe-debuginfod)

# Agenda

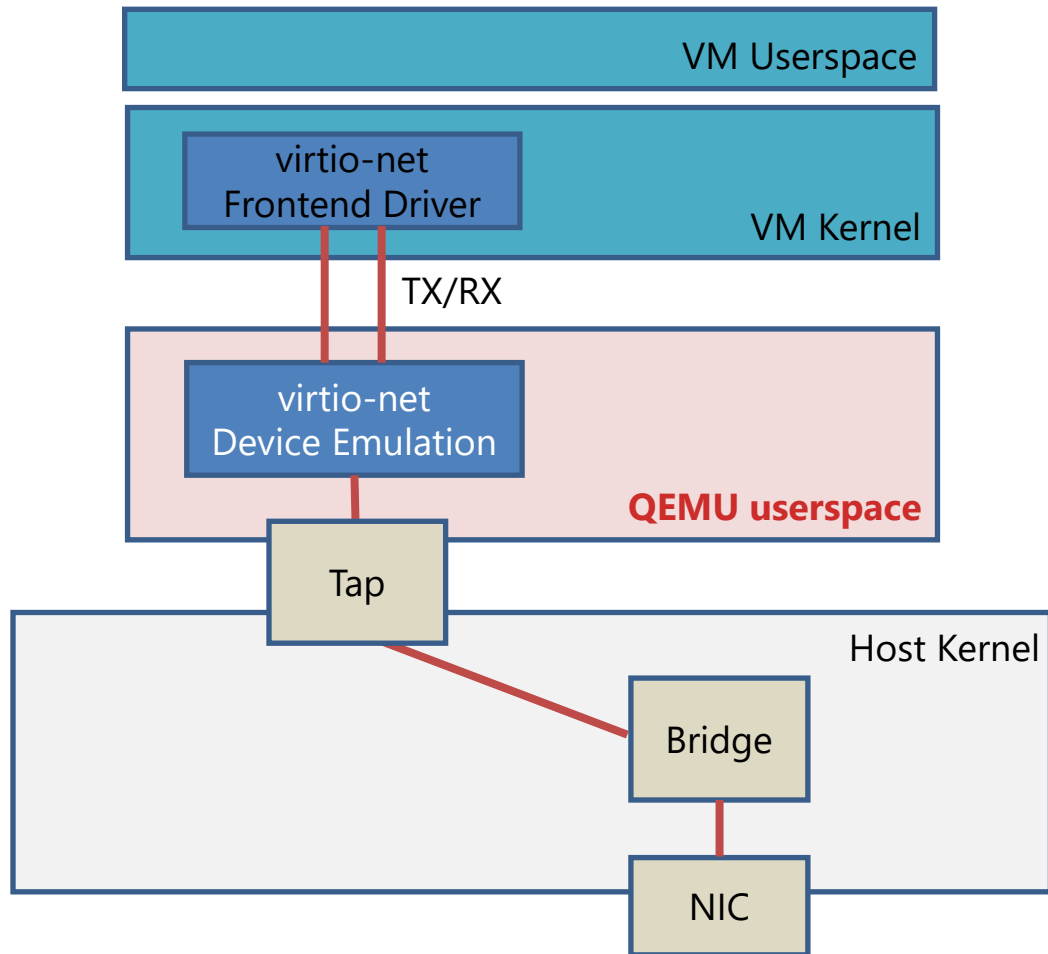
- Introduction
- Debugging Tools for Virtualized Systems
- **Example: Analyzing vhost-net**
- Summary

# Goal of This Section

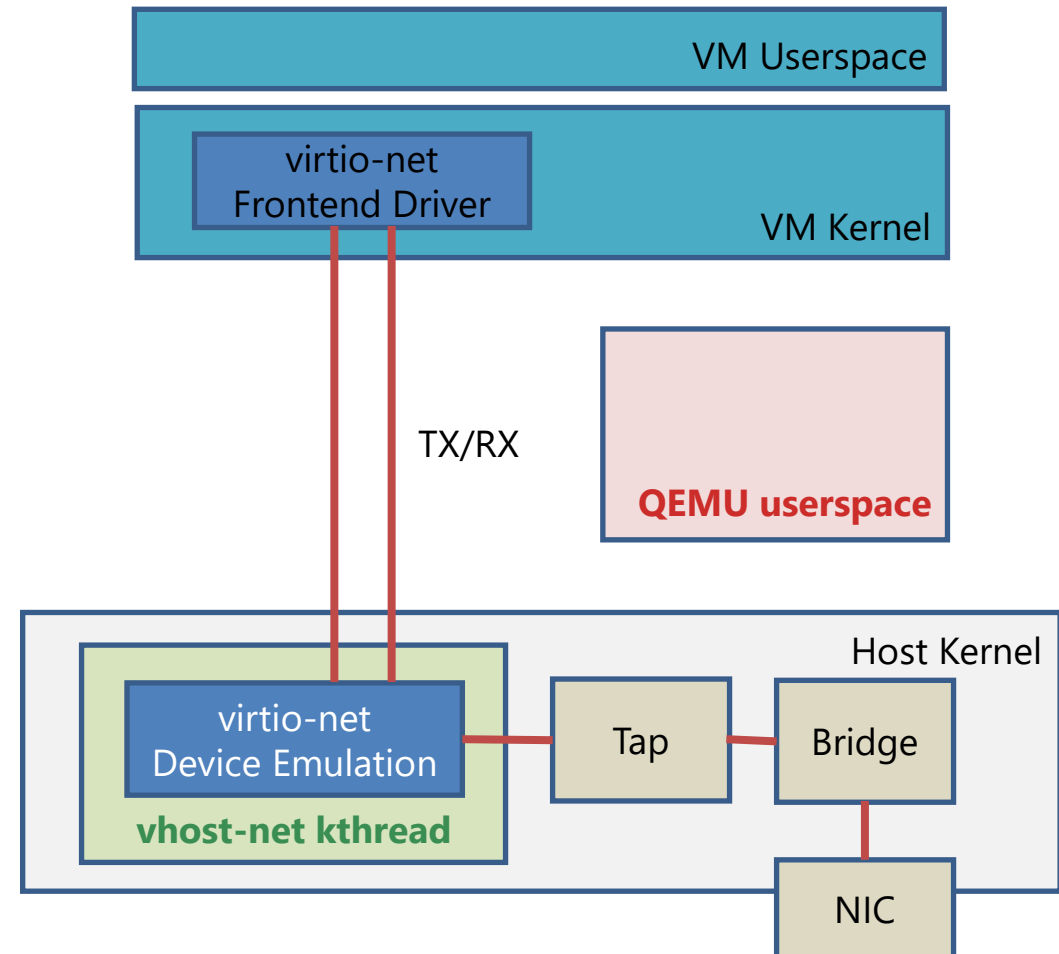
- Gaining knowledge of debugging tools' usage and techniques in a virtualized environment, through a practical example
- Investigating the behavior of **vhost-net** (acceleration for virtio-net), by comparing it with **standard virtio-net**
  - How and why it can improve performance?
  - Any potential side effects?

# Standard virtio-net vs vhost-net

## Standard virtio:



## vhost-net:




When **vhost-net** is enabled, **virtio-net** device emulation is shifted from **QEMU userspace** to the **vhost-net kernel thread**

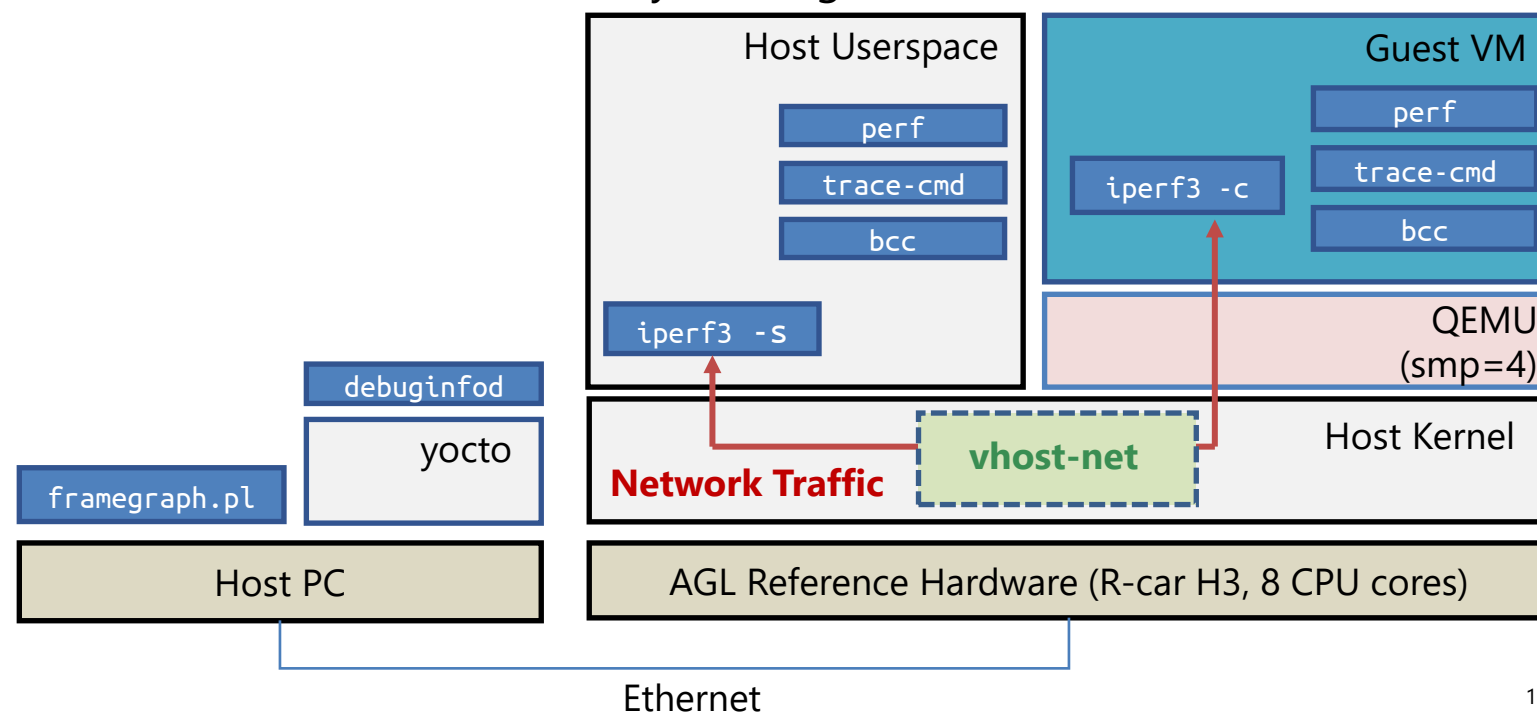


# Preconditions

- Base image: [agl-kvm-demo-platform](#) built from AGL-UCB master branch
- Additional installations: vhost-net, iperf3, ssh and various debugging tools
  - For more details, please refer to the setup procedures in the appendix
- Running on AGL Reference Hardware (R-car H3)
- Hosting a single guest VM by QEMU using smp=4 (4 vCPUs)
- Generating network traffic by iperf3 between host and guest
- Halting all workloads other than iperf3 prior to the investigation
- Comparing metrics/traces between **vhost-net** and **standard virtio** by utilizing **perf**, **trace-cmd**, and **bcc**



AGL Reference Hardware



# Measuring Network Bitrate: `iperf`

## Standard virtio:

```
root@guest:~# iperf3 -c 172.16.10.1
Connecting to host 172.16.10.1, port 5201
[ 5] local 172.16.10.2 port 42458 connected to 172.16.10.1 port 5201
[ ID] Interval           Transfer    Bitrate      Retr  Cwnd
[ 5]  0.00-1.00   sec    175 MBytes  1.47 Gbits/sec    0   1.53 MBytes
[ 5]  1.00-2.00   sec    191 MBytes  1.61 Gbits/sec    0   1.53 MBytes
[ 5]  2.00-3.00   sec    190 MBytes  1.59 Gbits/sec    0   1.61 MBytes
[ 5]  3.00-4.00   sec    192 MBytes  1.61 Gbits/sec    0   1.69 MBytes
[ 5]  4.00-5.00   sec    192 MBytes  1.61 Gbits/sec    0   1.80 MBytes
[ 5]  5.00-6.00   sec    210 MBytes  1.76 Gbits/sec    0   1.80 MBytes
[ 5]  6.00-7.00   sec    199 MBytes  1.67 Gbits/sec    0   1.80 MBytes
[ 5]  7.00-8.00   sec    202 MBytes  1.70 Gbits/sec    0   1.89 MBytes
[ 5]  8.00-9.00   sec    212 MBytes  1.78 Gbits/sec    0   1.99 MBytes
[ 5]  9.00-10.00  sec    215 MBytes  1.80 Gbits/sec    0   2.10 MBytes
- - - - -
[ ID] Interval           Transfer    Bitrate      Retr
[ 5]  0.00-10.00  sec    1.93 GBytes  1.66 Gbits/sec    0      sender
[ 5]  0.00-10.01  sec    1.93 GBytes  1.66 Gbits/sec      receiver
```

## vhost-net:

```
root@guest:~# iperf3 -c 172.16.10.1
Connecting to host 172.16.10.1, port 5201
[ 5] local 172.16.10.2 port 55840 connected to 172.16.10.1 port 5201
[ ID] Interval           Transfer    Bitrate      Retr  Cwnd
[ 5]  0.00-1.00   sec    391 MBytes  3.28 Gbits/sec    0   1.64 MBytes
[ 5]  1.00-2.00   sec    382 MBytes  3.21 Gbits/sec    0   1.64 MBytes
[ 5]  2.00-3.00   sec    385 MBytes  3.23 Gbits/sec    0   1.74 MBytes
[ 5]  3.00-4.00   sec    385 MBytes  3.23 Gbits/sec    0   1.74 MBytes
[ 5]  4.00-5.00   sec    382 MBytes  3.21 Gbits/sec    0   1.74 MBytes
[ 5]  5.00-6.00   sec    386 MBytes  3.23 Gbits/sec    0   1.74 MBytes
[ 5]  6.00-7.00   sec    401 MBytes  3.37 Gbits/sec    0   2.06 MBytes
[ 5]  7.00-8.00   sec    391 MBytes  3.29 Gbits/sec    0   2.06 MBytes
[ 5]  8.00-9.00   sec    386 MBytes  3.24 Gbits/sec    0   2.06 MBytes
[ 5]  9.00-10.00  sec    392 MBytes  3.29 Gbits/sec    0   2.06 MBytes
- - - - -
[ ID] Interval           Transfer    Bitrate      Retr
[ 5]  0.00-10.00  sec    3.79 GBytes  3.26 Gbits/sec    0      sender
[ 5]  0.00-10.01  sec    3.79 GBytes  3.26 Gbits/sec      receiver
```

vhost-net is about 2x faster than standard virtio-net

# KVM Event Stats: `perf stat -e 'kvm:*`

## Standard virtio:

```
root@host:~# perf stat -e 'kvm:*' -a -- ssh guest iperf3 -c 172.16.10.1 -n 100M
```

```
3656    kvm:vgic_update_irq_pending
1145    kvm:kvm_wfx_arm64
5453    kvm:kvm_arm_setup_debug
5453    kvm:kvm_arm_clear_debug
10906   kvm:kvm_arm_set_dreg32
2436    kvm:kvm_handle_sys_reg
2436    kvm:kvm_sys_access
5453    kvm:kvm_entry
5453    kvm:kvm_exit
1613    kvm:kvm_guest_fault
1181    kvm:kvm_irq_line
2475    kvm:kvm_timer_update_irq
7511    kvm:kvm_get_timer_map
2328    kvm:kvm_timer_save_state
2328    kvm:kvm_timer_restore_state
2328    kvm:kvm_timer_emulate
1180    kvm:kvm_userspace_exit
1143    kvm:kvm_vcpu_wakeup
2201    kvm:kvm_mmio
159     kvm:kvm_halt_poll_ns
```

Limiting number of packets

## vhost-net:

```
root@host:~# perf stat -e 'kvm:*' -a -- ssh guest iperf3 -c 172.16.10.1 -n 100M
```

```
5377    kvm:vgic_update_irq_pending
611     kvm:kvm_wfx_arm64
7348    kvm:kvm_arm_setup_debug
7348    kvm:kvm_arm_clear_debug
14696   kvm:kvm_arm_set_dreg32
2429    kvm:kvm_handle_sys_reg
2429    kvm:kvm_sys_access
7348    kvm:kvm_entry
7348    kvm:kvm_exit
3510    kvm:kvm_guest_fault
2548    kvm:kvm_irq_line
2829    kvm:kvm_timer_update_irq
12315   kvm:kvm_get_timer_map
2678    kvm:kvm_timer_save_state
2678    kvm:kvm_timer_restore_state
2678    kvm:kvm_timer_emulate
2188    kvm:kvm_userspace_exit
611     kvm:kvm_vcpu_wakeup
4604    kvm:kvm_mmio
128     kvm:kvm_halt_poll_ns
```

Significant reduction in WFXx (CPU Standby) and wakeup occurrences on guest vCPUs

# KVM VM-Exit Stats: `perf kvm stat`

## Standard virtio:

```
root@host:~# perf kvm stat record -a -- ssh guest iperf3 -c 172.16.10.1 -n 500M && perf kvm report
```

Analyze events for all VMs, all VCPUs:

VM-EXIT	Samples	Samples%	Time%	Min Time	Max Time	Avg time
DABT_LOW	4578	40.12%	26.53%	6.60us	1961.41us	104.75us ( +- 1.95% )
WFX	3546	31.08%	46.04%	8.04us	2113.93us	234.70us ( +- 2.56% )
SYS64	2430	21.30%	2.18%	5.64us	1621.33us	16.18us ( +- 7.93% )
IRQ	856	7.50%	25.25%	11.16us	2793.49us	533.31us ( +- 3.70% )

## vhost-net:

```
root@host:~# perf kvm stat record -a -- ssh guest iperf3 -c 172.16.10.1 -n 500M && perf kvm report
```

Analyze events for all VMs, all VCPUs:

VM-EXIT	Samples	Samples%	Time%	Min Time	Max Time	Avg time
DABT_LOW	23050	66.71%	66.59%	6.24us	3449.53us	59.70us ( +- 1.13% )
IRQ	6750	19.54%	19.44%	8.28us	3951.14us	59.50us ( +- 4.89% )
SYS64	2433	7.04%	2.36%	5.76us	3352.33us	20.03us ( +- 10.95% )
WFX	2318	6.71%	11.62%	12.72us	1807.21us	103.55us ( +- 2.08% )

- While we've observed a decrease in WFX again, VM-Exit themselves have increased (mainly due to Data abort and IRQ events)
- Common belief: a higher number of VM-Exit events leads to poorer performance
  - But this may not be always true, WFX could be more costly

# Monitoring I/O Operations on Virtio: `virtiostat` (bcc)

## Standard virtio:

```
root@guest:~# python3 /usr/share/bcc/tools/virtiostat
[...]
      Driver   Device   VQ Name   In SGs Out SGs           In BW           Out BW
b'virtio_net' b'virtio0' b'input.0'   8161      0      12733952           0
b'virtio_net' b'virtio0' b'output.0'      0    8556           0      376253940
[...]
```

## vhost-net:

```
root@guest:~# python3 /usr/share/bcc/tools/virtiostat
[...]
      Driver   Device   VQ Name   In SGs Out SGs           In BW           Out BW
b'virtio_net' b'virtio0' b'input.0'   7385      0      11523072           0
b'virtio_net' b'virtio0' b'output.0'      0    8572           0      375796680
[...]
```

- Traffic handled by virtio-net is almost consistent
  - Behavior of the guest side should not be changed

# Performance Sampling Across Host and Guest: `perf kvm record`

Before using `perf kvm record`, it's necessary to mount the guest filesystem from the host, as shown below:

```
root@host:~# mntpoint="/tmp/guestmount/$(pgrep -f qemu-system-aarch64)"
root@host:~# mkdir -p "$mntpoint"
root@host:~# sshfs -o allow_other,direct_io guest:/"$mntpoint"
```

Afterward, we can pass the mountpoint using `--guestmount` option:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount record ...
```

# Performance Sampling Across Host and Guest: `perf kvm record`

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount record -F99 -a -- ssh guest iperf3 -c 172.16.10.1 -b 100M -t 1
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio --show-cpu-utilization
```

#	Overhead	sys	usr	guest sys	guest usr	Command	Shared Object	Symbol
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	20.11%	0.00%	0.00%	20.11%	0.00%	:1035	[guest.kernel.kallsyms.1026]	[g] receive_buf
	13.47%	13.47%	0.00%	0.00%	0.00%	swapper	[kernel.kallsyms]	[k] cpuidle_enter_state
	9.56%	0.00%	9.56%	0.00%	0.00%	qemu-system-aar	qemu-system-aarch64	[.] replay_mutex_unlock
	4.97%	0.00%	0.00%	4.97%	0.00%	:1037	[guest.kernel.kallsyms.1026]	[g] cpu_do_idle
	4.01%	4.01%	0.00%	0.00%	0.00%	swapper	[kernel.kallsyms]	[k] _raw_spin_unlock_irq
	3.19%	0.00%	3.19%	0.00%	0.00%	ssh	ssh.openssh	[.] square
	2.57%	0.00%	0.00%	2.57%	0.00%	:1035	[guest.kernel.kallsyms.1026]	[g] unwind_next
	2.00%	0.00%	0.00%	0.00%	2.00%	:1035	[unknown]	[u] 0x000000557d660af8
	1.84%	0.00%	0.00%	1.84%	0.00%	:1037	[guest.kernel.kallsyms.1026]	[g] update_load_avg
	1.82%	0.00%	1.82%	0.00%	0.00%	qemu-system-aar	qemu-system-aarch64	[.] object_dynamic_cast_assert
	1.82%	0.00%	1.82%	0.00%	0.00%	qemu-system-aar	libc.so.6	[.] clock_gettime@@GLIBC_2.17
	1.78%	0.00%	0.00%	1.78%	0.00%	:1037	[guest.kernel.kallsyms.1026]	[g] unwind_next
	1.68%	1.68%	0.00%	0.00%	0.00%	qemu-system-aar	[kernel.kallsyms]	[k] _raw_spin_unlock_irq
	1.54%	0.00%	0.00%	1.54%	0.00%	:1037	[guest.kernel.kallsyms.1026]	[g] __schedule

guest kernel function  
named receive\_buf()  
consumes 20.11% of CPU

# Performance Sampling Across Host and Guest: `perf kvm record`

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount record -F99 -a -- ssh guest iperf3 -c 172.16.10.1 -b 100M -t 1
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio --show-cpu-utilization
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command	Stack	Symbol
#	.....	.....	.....	.....	.....	.....	.....	.....	.....	.....
#	20.11%	0.00%	0.00%	20.11%	0.00%	:	1035			receive_buf
	13.47%	13.47%	0.00%	0.00%	0.00%	swapper				cpuidle_enter_state
	9.56%	0.00%	9.56%	0.00%	0.00%	qemu-system-aar		qemu-system-aarch64	[.]	replay_mutex_unlock
	4.97%	0.00%	0.00%	4.97%	0.00%	:1037		[guest.kernel.kallsyms.1026]	[g]	cpu_do_idle
	4.01%	4.01%	0.00%	0.00%	0.00%	swapper		[kernel.kallsyms]	[k]	_raw_spin_unlock_irq
	3.19%	0.00%	3.19%	0.00%	0.00%	ssh		ssh.openssh	[.]	square
	2.57%	0.00%	0.00%	2.57%	0.00%	:1035		[guest.kernel.kallsyms.1026]	[g]	unwind_next
	2.00%	0.00%	0.00%	0.00%	2.00%	:1035		[unknown]	[u]	0x000000557d660af8
	1.84%	0.00%	0.00%	1.84%	0.00%	:1037		[guest.kernel.kallsyms.1026]	[g]	update_load_avg
	1.82%	0.00%	1.82%	0.00%	0.00%	qemu-system-aar		qemu-system-aarch64	[.]	object_dynamic_cast_assert
	1.82%	0.00%	1.82%	0.00%	0.00%	qemu-system-aar		libc.so.6	[.]	clock_gettime@@GLIBC_2.17
	1.78%	0.00%	0.00%	1.78%	0.00%	:1037		[guest.kernel.kallsyms.1026]	[g]	unwind_next
	1.68%	1.68%	0.00%	0.00%	0.00%	qemu-system-aar		[kernel.kallsyms]	[k]	_raw_spin_unlock_irq
	1.54%	0.00%	0.00%	1.54%	0.00%	:1037		[guest.kernel.kallsyms.1026]	[g]	__schedule

**:tid** represents a child thread of another process (in this case, vCPU threads of qemu)



# Performance Sampling Across Host and Guest: `perf kvm record`

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount record -F99 -a -- ssh guest iperf3 -c 172.16.10.1 -b 100M -t 1
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio --show-cpu-utilization
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command	Shared Object	Symbol
#	.....	.....	.....	.....	.....	.....	.....	.....	.....	.....
#										
	20.11%	0.00%	0.00%	20.11%	0.00%			:1035	[guest.kernel.kallsyms.1026]	[g] receive_buf
	13.47%	13.47%	0.00%	0.00%	0.00%			swapper	[kernel.kallsyms]	[k] cpuidle_enter_state
	9.56%	0.00%	9.56%	0.00%	0.00%			qemu-system-aar	qemu-system-aarch64	[.] replay_mutex_unlock
	4.97%	0.00%	0.00%	4.97%	0.00%			:1037	[guest.kernel.kallsyms.1026]	[g] cpu_do_idle
	4.01%	4.01%	0.00%	0.00%	0.00%			swapper	[kernel.kallsyms]	[k] _raw_spin_unlock_irq
	3.19%	0.00%	3.19%	0.00%	0.00%			ssh	ssh.openssh	[.] square
	2.57%	0.00%	0.00%	2.57%	0.00%			:1035	[guest.kernel.kallsyms.1026]	[g] unwind_next
	2.00%	0.00%	0.00%	0.00%	2.00%			:1035	[unknown]	[u] 0x000000557d660af8
	1.84%	0.00%	0.00%	1.84%	0.00%			:1037	[guest.kernel.kallsyms.1026]	[g] update_load_avg
	1.82%	0.00%	1.82%	0.00%	0.00%			qemu-system-aar	qemu-system-aarch64	[.] object_dynamic_cast_assert
	1.82%	0.00%	1.82%	0.00%	0.00%			qemu-system-aar	libc.so.6	[.] clock_gettime@@GLIBC_2.17
	1.78%	0.00%	0.00%	1.78%	0.00%			:1037	[guest.kernel.kallsyms.1026]	[g] unwind_next
	1.68%	1.68%	0.00%	0.00%	0.00%			qemu-system-aar	[kernel.kallsyms]	[k] _raw_spin_unlock_irq
	1.54%	0.00%	0.00%	1.54%	0.00%			:1037	[guest.kernel.kallsyms.1026]	[g] __schedule

- Can analyze high-load functions across both host and guest systems
- However, this function-level of detail might be too much in this analysis

# Performance Sampling Across Host and Guest: `perf kvm record` **Panasonic** AUTOMOTIVE

- Summarize the results using `perf kvm report --sort` option

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -  
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	30.72%	0.00%	0.00%	28.72%	2.00%			:1035
	25.31%	4.38%	20.94%	0.00%	0.00%			qemu-system-aar
	18.23%	0.00%	0.00%	18.23%	0.00%			:1037
	18.13%	18.13%	0.00%	0.00%	0.00%			swapper
	3.98%	0.00%	3.98%	0.00%	0.00%			ssh
	1.35%	0.00%	0.00%	1.35%	0.00%			:1038
	1.29%	0.85%	0.44%	0.00%	0.00%			iperf3
	0.97%	0.00%	0.00%	0.97%	0.00%			:1039
	0.01%	0.01%	0.01%	0.00%	0.00%			perf
	0.01%	0.01%	0.00%	0.00%	0.00%			migration/6
	0.01%	0.01%	0.00%	0.00%	0.00%			migration/5

## vhost-net:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -  
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	40.37%	0.00%	0.00%	40.37%	0.00%			:1146
	15.19%	15.19%	0.00%	0.00%	0.00%			swapper
	13.08%	13.08%	0.00%	0.00%	0.00%			qemu-system-aar
	11.72%	0.00%	0.00%	11.72%	0.00%			:1147
	6.87%	0.00%	6.87%	0.00%	0.00%			ssh
	5.45%	0.00%	0.00%	5.45%	0.00%			:1149
	4.28%	4.28%	0.00%	0.00%	0.00%			iperf3
	2.78%	2.78%	0.00%	0.00%	0.00%			vhost-1134
	0.21%	0.00%	0.21%	0.00%	0.00%			sshd
	0.01%	0.01%	0.01%	0.00%	0.00%			perf
	0.01%	0.01%	0.00%	0.00%	0.00%			migration/0
	0.01%	0.01%	0.00%	0.00%	0.00%			migration/5

# Performance Sampling Across Host and Guest: `perf kvm record` **Panasonic** AUTOMOTIVE

- Summarize the results using `perf kvm report --sort` option

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -  
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	30.72%	0.00%	0.00%	28.72%	2.00%	:	1035	
	25.31%	4.38%	20.94%	0.00%	0.00%	qemu-system-aar		
	18.23%	0.00%	0.00%	18.23%	0.00%	:	1037	
	18.13%	18.13%	0.00%	0.00%	0.00%	swapper		
	3.98%	0.00%	3.98%	0.00%	0.00%	ssh		
	1.35%	0.00%	0.00%	1.35%	0.00%	:	1038	
	1.29%	0.85%	0.44%	0.00%	0.00%	iperf3		
	0.97%	0.00%	0.00%	0.97%	0.00%	:	1039	
	0.01%	0.01%	0.01%	0.00%	0.00%	perf		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/6		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5		

These overheads  
includes  
idle functions

## vhost-net:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -  
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	40.37%	0.00%	0.00%	40.37%	0.00%	:	1146	
	15.19%	15.19%	0.00%	0.00%	0.00%	swapper		
	13.08%	13.08%	0.00%	0.00%	0.00%	qemu-system-aar		
	11.72%	0.00%	0.00%	11.72%	0.00%	:	1147	
	6.87%	0.00%	6.87%	0.00%	0.00%	ssh		
	5.45%	0.00%	0.00%	5.45%	0.00%	:	1149	
	4.28%	4.28%	0.00%	0.00%	0.00%	iperf3		
	2.78%	2.78%	0.00%	0.00%	0.00%	vhost-1134		
	0.21%	0.00%	0.21%	0.00%	0.00%	sshd		
	0.01%	0.01%	0.01%	0.00%	0.00%	perf		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/0		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5		

# Performance Sampling Across Host and Guest: `perf kvm record` **Panasonic** AUTOMOTIVE

- Summarize the results using `perf kvm report --sort` option

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest sys	guest usr	Command
#	.....	.....	.....	.....	.....	.....
#						
	30.72%	0.00%	0.00%	28.72%	2.00%	:1035
	25.31%	4.38%	20.94%	0.00%	0.00%	qemu-system-aar
	18.23%	0.00%	0.00%	18.23%	0.00%	:1037
	18.13%	18.13%	0.00%	0.00%	0.00%	swapper
	3.98%	0.00%	3.98%	0.00%	0.00%	ssh
	1.35%	0.00%	0.00%	1.35%	0.00%	:1038
	1.29%	0.85%	0.44%	0.00%	0.00%	iperf3
	0.97%	0.00%	0.00%	0.97%	0.00%	:1039
	0.01%	0.01%	0.01%	0.00%	0.00%	perf
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/6
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization | grep idle | sed -e 's!\[.*\]!!'
```

13.47%	13.47%	0.00%	0.00%	0.00%	swapper	cpuidle_enter_state
4.97%	0.00%	0.00%	4.97%	0.00%	:1037	cpu_do_idle
0.64%	0.64%	0.00%	0.00%	0.00%	swapper	tick_nohz_idle_exit
0.53%	0.00%	0.00%	0.53%	0.00%	:1035	cpu_do_idle

## vhost-net:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest sys	guest usr	Command
#	.....	.....	.....	.....	.....	.....
#						
	40.37%	0.00%	0.00%	40.37%	0.00%	:1146
	15.19%	15.19%	0.00%	0.00%	0.00%	swapper
	13.08%	13.08%	0.00%	0.00%	0.00%	qemu-system-aar
	11.72%	0.00%	0.00%	11.72%	0.00%	:1147
	6.87%	0.00%	6.87%	0.00%	0.00%	ssh
	5.45%	0.00%	0.00%	5.45%	0.00%	:1149
	4.28%	4.28%	0.00%	0.00%	0.00%	iperf3
	2.78%	2.78%	0.00%	0.00%	0.00%	vhost-1134
	0.21%	0.00%	0.21%	0.00%	0.00%	sshd
	0.01%	0.01%	0.01%	0.00%	0.00%	perf
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/0
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization | grep idle | sed -e 's!\[.*\]!!'
```

30.01%	0.00%	0.00%	30.01%	0.00%	:1146	cpu_do_idle
13.05%	13.05%	0.00%	0.00%	0.00%	swapper	cpuidle_enter_state
6.55%	0.00%	0.00%	6.55%	0.00%	:1147	cpu_do_idle
3.59%	0.00%	0.00%	3.59%	0.00%	:1149	cpu_do_idle

identify the idle functions and manually exclude them from each processes

# Performance Sampling Across Host and Guest: perf kvm record



- Summarize the results using `perf kvm report --sort` option

## Standard virtio:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	30.72%	0.00%	0.00%	28.72%	2.00%	:	1035	
	25.31%	4.38%	20.94%	0.00%	0.00%	qemu-system-aar		
	18.23%	0.00%	0.00%	18.23%	0.00%	:	1037	
	18.13%	18.13%	0.00%	0.00%	0.00%	swapper		
	3.98%	0.00%	3.98%	0.00%	0.00%	ssh		
	1.35%	0.00%	0.00%	1.35%	0.00%	:	1038	
	1.29%	0.85%	0.44%	0.00%	0.00%	iperf3		
	0.97%	0.00%	0.00%	0.97%	0.00%	:	1039	
	0.01%	0.01%	0.01%	0.00%	0.00%	perf		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/6		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5		

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization | grep idle | sed -e
```

13.47%	13.47%	0.00%	0.00%
4.97%	0.00%	0.00%	4.97%
0.64%	0.64%	0.00%	0.00%
0.53%	0.00%	0.00%	0.53%

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
	30.19%	0.00%	0.00%	28.19%	2.00%	:	1035	
	25.31%	4.38%	20.94%	0.00%	0.00%	qemu-system-aar		
	13.26%	0.00%	0.00%	13.26%	0.00%	:	1037	

## vhost-net:

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization --sort=comm
```

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
#								
	40.37%	0.00%	0.00%	40.37%	0.00%	:	1146	
	15.19%	15.19%	0.00%	0.00%	0.00%	swapper		
	13.08%	13.08%	0.00%	0.00%	0.00%	qemu-system-aar		
	11.72%	0.00%	0.00%	11.72%	0.00%	:	1147	
	6.87%	0.00%	6.87%	0.00%	0.00%	ssh		
	5.45%	0.00%	0.00%	5.45%	0.00%	:	1149	
	4.28%	4.28%	0.00%	0.00%	0.00%	iperf3		
	2.78%	2.78%	0.00%	0.00%	0.00%	vhost-1134		
	0.21%	0.00%	0.21%	0.00%	0.00%	sshd		
	0.01%	0.01%	0.01%	0.00%	0.00%	perf		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/0		
	0.01%	0.01%	0.00%	0.00%	0.00%	migration/5		

```
root@host:~# perf kvm --host --guest --guestmount=/tmp/guestmount report --stdio -
show-cpu-utilization | grep idle | sed -e 's!\[.*\]!!'
```

0.00%	0.00%	30.01%	0.00%	:	1146	cpu_do_idle
13.05%	0.00%	0.00%	0.00%	:	swapper	cpuidle_enter_state
0.00%	0.00%	6.55%	0.00%	:	1147	cpu_do_idle
3.59%	0.00%	3.59%	0.00%	:	1149	cpu_do_idle

#	Overhead	sys	usr	guest	sys	guest	usr	Command
#	.....	.....	.....	.....	.....	.....	.....	.....
	10.36%	0.00%	0.00%	10.36%	0.00%	:	1146	
	13.08%	13.08%	0.00%	0.00%	0.00%	qemu-system-aar		
	5.17%	0.00%	0.00%	5.17%	0.00%	:	1147	
	1.86%	0.00%	0.00%	1.86%	0.00%	:	1149	

Results after exclusion:  
CPU consumption has significantly decreased  
in vhost-net

# Workload Visualization: `framegraph.pl`

## Basic usage:

1. Perform a `perf record` with stack-trace (`-g`) and convert the data into text format using `perf script`

```
# perf record -a -g -F99 -- sleep 1    # or, any command you want to profile
# perf script > script.txt
```

2. Transfer the output to your host PC, then convert it to a SVG file using `stackcollapse-perf.pl` and `framegraph.pl`

```
$ cat script.txt | stackcollapse-perf.pl | framegraph.pl > output.svg
```

3. Open the SVG file in a web browser to view the results.



# Workload Visualization: `framegraph.pl`

## Tips & Tricks for virtualized systems:

1. Record samples on both host and guest systems at once

The diagram shows a terminal command with three annotations in yellow boxes: "Record on host side" pointing to the host perf record command, "Record on guest side" pointing to the guest perf record command, and "Workload" pointing to the iperf3 command. The command is: `root@host:~# perf record -a -g -F99 -o /tmp/perf.data.host -- ssh guest perf record -a -g -F99 -o /tmp/perf.data.guest \ iperf3 -c 172.16.10.1 -b 1G`

2. Merge the two outputs into a single SVG file


```
$ sed -i -e 's!\!(^\^[^ \t]\+\\)\!<guest>\!1!' script.guest # Add prefix to guest output
$ cat script.host script.guest | stackcollapse-perf.pl | flamegraph.pl > output.svg
```

\*Note that these procedures are designed for convenience and may not be strictly correct.

There may be overlapping workloads or inconsistent scaling between the guest and host systems.\*

**Panasonic**  
AUTOMOTIVE

## vhost-net:




- 32



# Workload Visualization: `framegraph.pl`


## Standard virtio:



<guest>iperf3

<guest>swapper

## vhost-net:




<guest>iperf3

<guest>swapper

Guest side: Similar workload between **standard virtio** and **vhost-net**


# Workload Visualization: framegraph.pl

## Standard virtio:



iperf3

## vhost-net:




iperf3

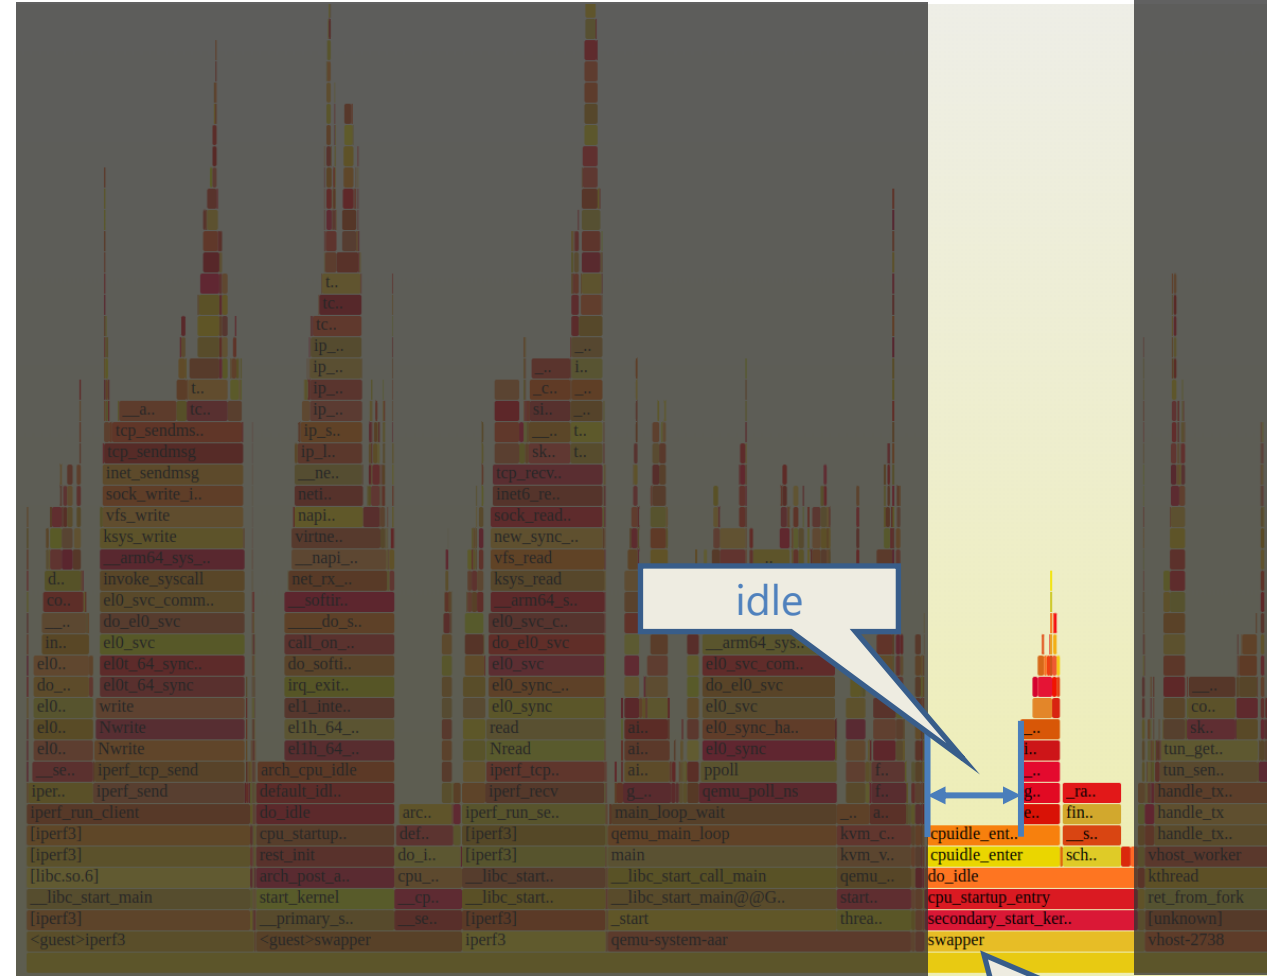
iperf3 on host: Similar workload between standard virtio and vhost-net

# Workload Visualization: framegraph.pl

## Standard virtio:




## vhost-net:



Swapper on host: Somewhat similar, but **idle-time** is slightly longer in **vhost-net**


# Workload Visualization: framegraph.pl

## Standard virtio:



qemu

## vhost-net:



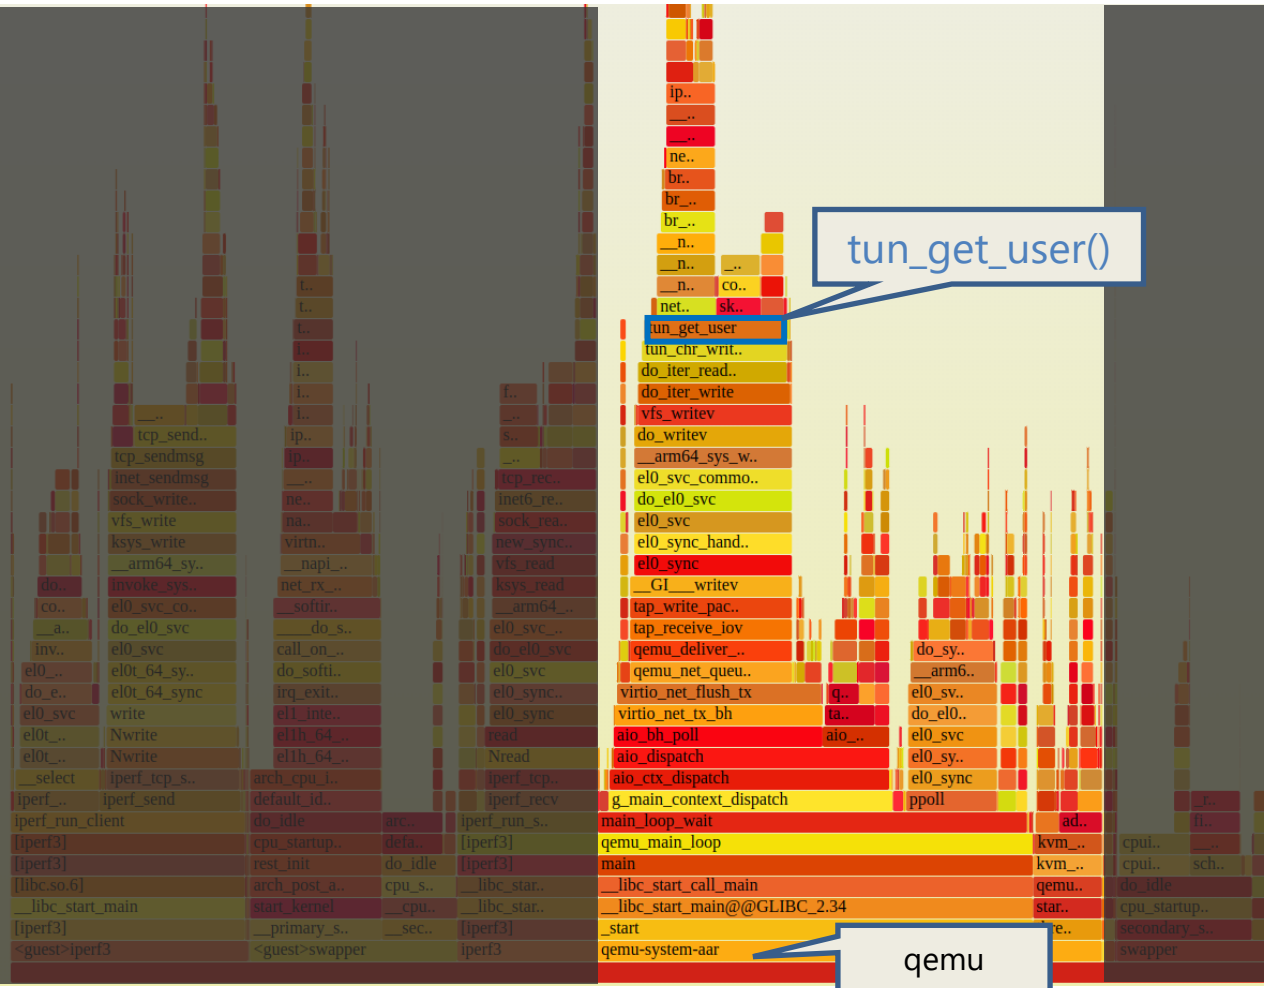
qemu

vhost


Primary difference lies in qemu and vhost kthread

# Workload Visualization: framegraph.pl

## Standard virtio:




## vhost-net:




- Same function load (tun\_get\_user) appears in different contexts (qemu vs vhost)
  - Visualize the shift of network device emulation

# Workload Visualization: framegraph.pl

## Standard virtio:



## vhost-net:




- **Standard virtio:** As qemu is a user process, it must utilize system calls (`writew`) to invoke kernel functions (`tun_get_user`)
- **vhost-net:** There is **no system calls** because vhost is a kernel thread




# Workload Visualization: `framegraph.pl`

## Standard virtio:




## vhost-net:

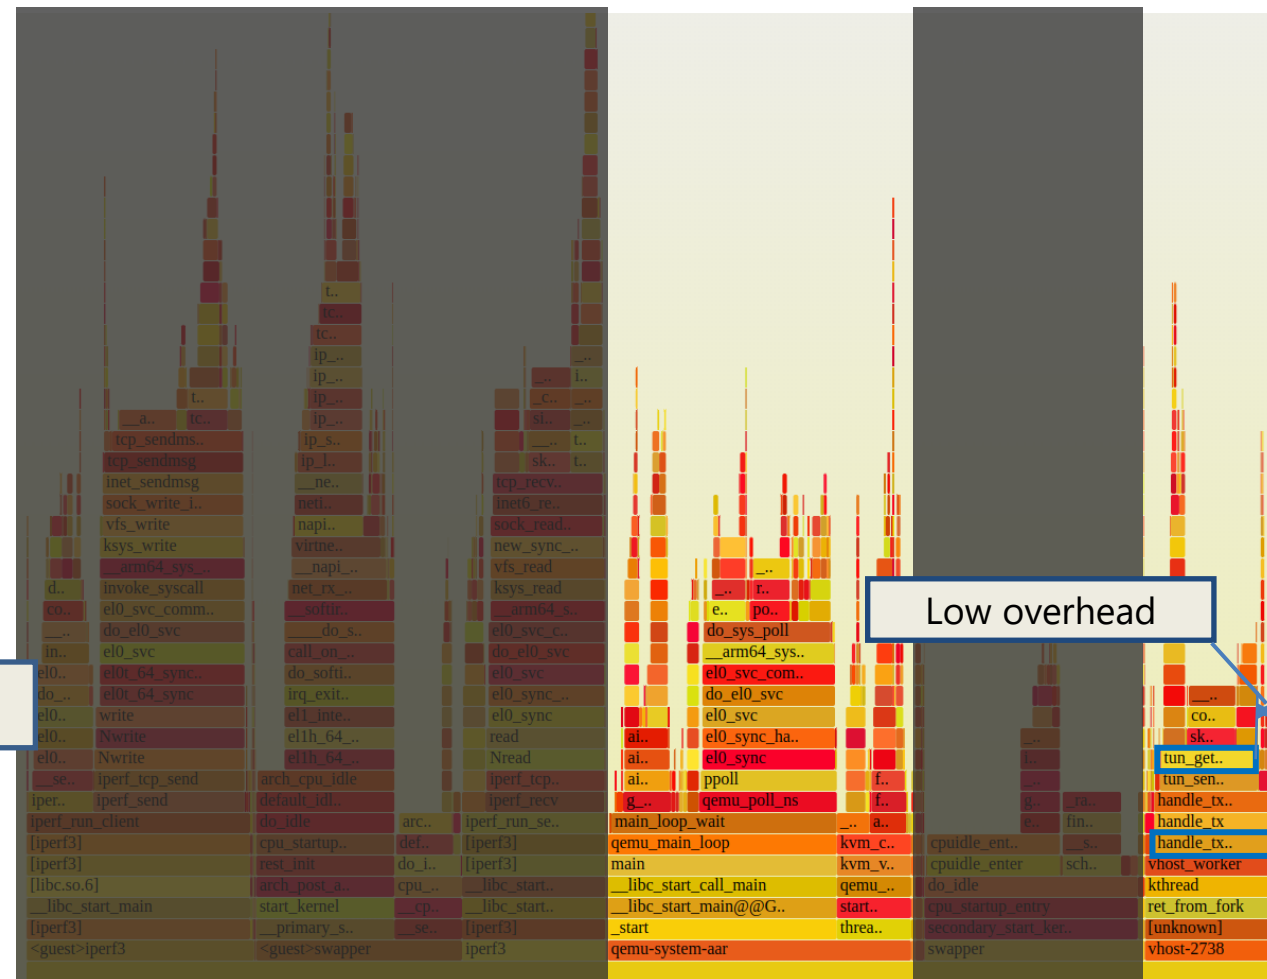


# Workload Visualization: framegraph.pl

## Standard virtio:



## vhost-net:



qemu generates **significant overhead** during network packet transfer,  
primarily due to **use of system calls**



# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

- To effectively analyze `context switches between host and guest`, I recommend tracing both '`kvm`' and '`sched`' events at once (`trace-cmd -e kvm -e sched`)
- Additionally, `kprobes` can be utilized to `trace almost all kernel functions`

Example: Creating and tracing a kprobe event for `tun_get_user`:

```
# Ensure the target function is traceable
root@host:~# grep tun_get_user /sys/kernel/debug/tracing/available_filter_functions
tun_get_user

# Create a kprobe event for tun_get_user
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events

# Verify that the kprobe event was successfully created
root@host:~# ls /sys/kernel/debug/tracing/events/kprobes/kp_tun_get_user
enable filter format id trigger

# Trace it with trace-cmd
root@host:~# trace-cmd record -e 'kprobes:kp_tun_get_user' -- sleep 1
```

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:        comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:       qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
      <idle>-0 [007] 7593.923936: sched_switch:          swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:       guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:    comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:          qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:    CTL: 0x000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:    (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:        Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:        comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:       qemu-system-aar:1361 [120] success=1 CPU:006
      <idle>-0 [006] 7593.924402: sched_switch:          swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq:    VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:    arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:       vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:      wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault: ipa 0xa003000, hsr 0x93810046, hxfar 0xffffffffc009675a50, pc 0xffffffffc00872cd74
(DABT_LOW), PC: 0xffffffffc00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio: mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking: comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup: qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry: PC: 0xffffffffc00872cd78
<idle>-0 [007] 7593.923936: sched_switch: swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit: <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xffffffffc008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64: guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime: comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch: qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map: CTL: 0x000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user: (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line: Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking: comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup: qemu-system-aar:1361 [120] success=1 CPU:006
qemu-system-aar-1353 [006] 7593.924402: sched_switch: swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate: arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns: vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup: wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio: mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" > /dev/trace
root@host:~# trace-cmd record -e kvm -e sched -e mmio 72.16.10.1 -n 100M && trace-cmd report
```

mmio event initiated by guest prompts  
host to transfer the next packets

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit: <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xffffffc00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fpu: ipa 0xa003000, hsr 0x93810046, hxfar 0xffffffc009675a50, pc 0xffffffc00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio: mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking: comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup: qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry: PC: 0xffffffc00872cd78
<idle>-0 [007] 7593.923936: sched_switch: swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit: <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xffffffc008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64: guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime: comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch: qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state: CTL: 0x000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user: (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line: Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking: comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup: qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch: swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate: arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns: vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup: wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio: mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd74
<idle>-0 [007] 7593.923936: sched_switch:       swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:       guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:   comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:       qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:    CTL: 0x0000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:    (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:       Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:       comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:       swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:    arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:      vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:      wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

qemu I/O thread awakens

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:      comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:      swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:      guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:  runtime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user: (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line: Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:      comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
qemu-system-aar-1353 [006] 7593.924402: sched_switch:      swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map: VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate: arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns: vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup: wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

guest vCPU enters sleep state using WfX,  
awaiting completion of network transfer



# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:       swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:       guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:   comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:       qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:   CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:   (tun)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:       Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:       comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:       swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:   arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:    vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:    wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

qemu I/O thread invokes tun\_get\_user()

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:       swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WFX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:       guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:   comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:       qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:   CTL: 0x0000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:    (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:       Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:       swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:    VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL:          0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:   arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:    vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:    wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

qemu handles IRQ event for the guest



# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:      swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:      guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:  comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:      qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:  CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:   (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:      Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:       comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:      swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:   arch_timer_ctx_index (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:   vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:   wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

guest vCPU gets woken up from sleep state

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:      comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:      swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:      guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:   comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990: sched_switch:      qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state:   CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:   (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:      Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:      comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:      swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:   arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:   vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:   wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

guest invokes next mmio event

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## Standard virtio:

```
root@host:~# echo "p:kp_tun_get_user tun_get_user" >> /sys/kernel/debug/tracing/kprobe_events
```

```
root@host:~# trace-cmd record -e kvm -e sched -e 'kprobes:kp_tun_get_user' -- ssh guest iperf3 -c 172.16.10.1 -n 100M && trace-cmd report
```

```
qemu-system-aar-1361 [006] 7593.923907: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024
(DABT_LOW), PC: 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923908: kvm_guest_fault:   ipa 0xa003000, hsr 0x93810046, hxfar 0xfffffff009675a50, pc 0xfffffff00872cd74
qemu-system-aar-1361 [006] 7593.923910: kvm_mmio:          mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1361 [006] 7593.923919: sched_waking:       comm=qemu-system-aar pid=1353 prio=120 target_cpu=007
qemu-system-aar-1361 [006] 7593.923926: sched_wakeup:      qemu-system-aar:1353 [120] success=1 CPU:007
qemu-system-aar-1361 [006] 7593.923934: kvm_entry:          PC: 0xfffffff00872cd78
<idle>-0 [007] 7593.923936: sched_switch:      swapper/7:0 [120] R ==> qemu-system-aar:1353 [120]
qemu-system-aar-1361 [006] 7593.923971: kvm_exit:          <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0001 (WfX), PC: 0xfffffff008be07e8
qemu-system-aar-1361 [006] 7593.923972: kvm_wfx_arm64:      guest executed wf>c< at: 0x0000000000000000
qemu-system-aar-1361 [006] 7593.923977: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923986: sched_stat_runtime:  comm=qemu-system-aar pid=1361 runtime=587763 [ns] vruntime=111836102550 [ns]
qemu-system-aar-1361 [006] 7593.923990:          qemu-system-aar:1361 [120] S ==> swapper/6:0 [120]
qemu-system-aar-1361 [006] 7593.923994:          VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.923996: kvm_timer_save_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.923997: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1353 [007] 7593.924129: kp_tun_get_user:   (ffff800010b335d0)
qemu-system-aar-1353 [007] 7593.924377: kvm_irq_line:      Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1353 [007] 7593.924379: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1353 [007] 7593.924386: sched_waking:       comm=qemu-system-aar pid=1361 prio=120 target_cpu=006
qemu-system-aar-1353 [007] 7593.924392: sched_wakeup:      qemu-system-aar:1361 [120] success=1 CPU:006
<idle>-0 [006] 7593.924402: sched_switch:      swapper/6:0 [120] R ==> qemu-system-aar:1361 [120]
qemu-system-aar-1361 [006] 7593.924417: kvm_get_timer_map:   VCPU: 0, dv: 1, dp: -1, ep: 0
qemu-system-aar-1361 [006] 7593.924419: kvm_timer_update_irq: VCPU: 0, IRQ 27, level 0
qemu-system-aar-1361 [006] 7593.924420: vgic_update_irq_pending: VCPU: 0, IRQ 27, level: 0
qemu-system-aar-1361 [006] 7593.924431: kvm_timer_restore_state: CTL: 0x0000001 CVAL: 0xf6643c1 arch_timer_ctx_index: 1
qemu-system-aar-1361 [006] 7593.924432: kvm_timer_emulate:   arch_timer_ctx_index: 0 (should_fire: 0)
qemu-system-aar-1361 [006] 7593.924440: kvm_halt_poll_ns:   vcpu 0: halt_poll_ns 10000 (grow 0)
qemu-system-aar-1361 [006] 7593.924441: kvm_vcpu_wakeup:   wait time 462722 ns, polling valid
...
qemu-system-aar-1361 [006] 7593.924478: kvm_mmio:          mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

**+411us**

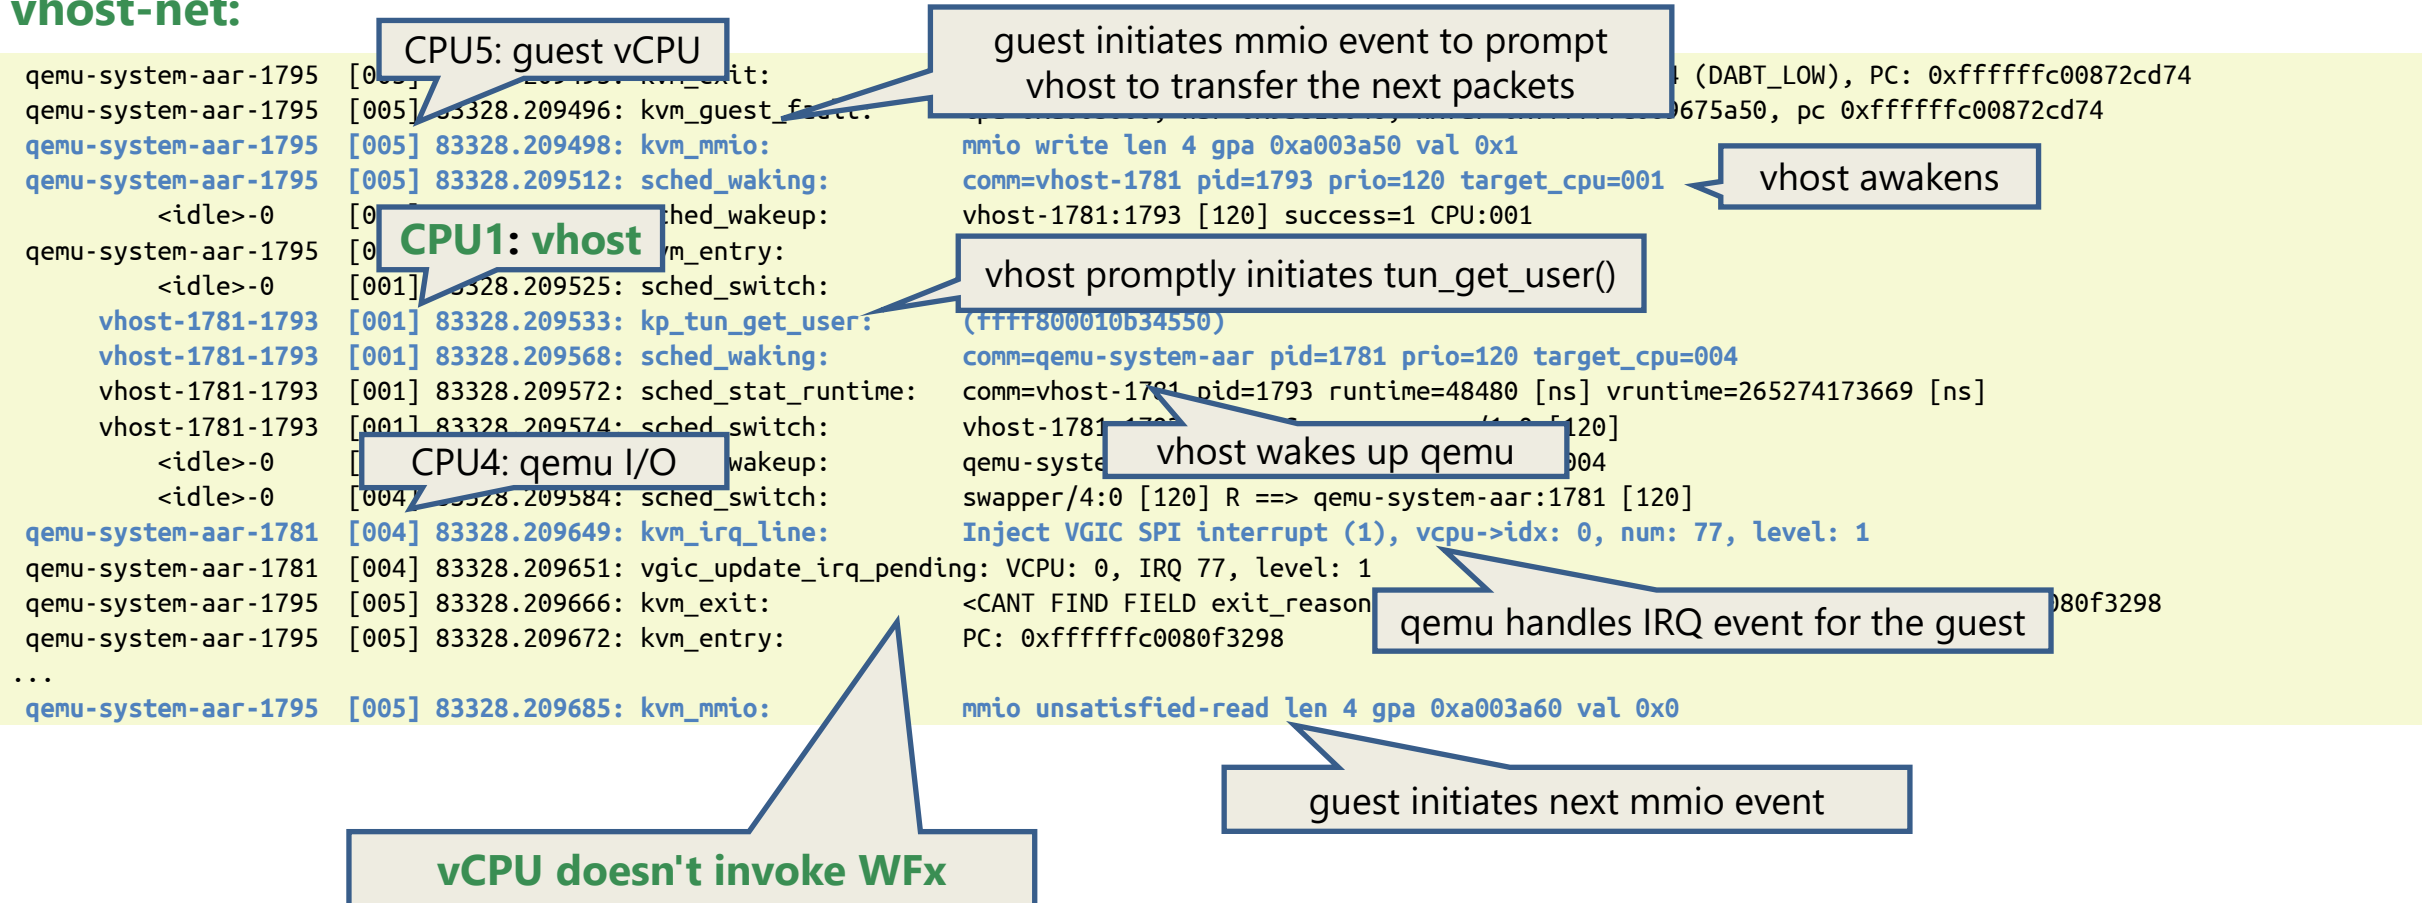
# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## vhost-net:

```
qemu-system-aar-1795 [005] 83328.209495: kvm_exit:      <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024 (DABT_LOW), PC: 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209496: kvm_guest_fault: ipa 0xa003000, hsr 0x93810046, hxfar 0xffffffffc009675a50, pc 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209498: kvm_mmio:      mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1795 [005] 83328.209512: sched_waking:    comm=vhost-1781 pid=1793 prio=120 target_cpu=001
<idle>-0 [001] 83328.209522: sched_wakeup:    vhost-1781:1793 [120] success=1 CPU:001
qemu-system-aar-1795 [005] 83328.209525: kvm_entry:      PC: 0xffffffffc00872cd78
<idle>-0 [001] 83328.209525: sched_switch:    swapper/1:0 [120] R ==> vhost-1781:1793 [120]
vhost-1781-1793 [001] 83328.209533: kp_tun_get_user:    (ffff800010b34550)
vhost-1781-1793 [001] 83328.209568: sched_waking:    comm=qemu-system-aar pid=1781 prio=120 target_cpu=004
vhost-1781-1793 [001] 83328.209572: sched_stat_runtime: comm=vhost-1781 pid=1793 runtime=48480 [ns] vruntime=265274173669 [ns]
vhost-1781-1793 [001] 83328.209574: sched_switch:    vhost-1781:1793 [120] S ==> swapper/1:0 [120]
<idle>-0 [004] 83328.209577: sched_wakeup:    qemu-system-aar:1781 [120] success=1 CPU:004
<idle>-0 [004] 83328.209584: sched_switch:    swapper/4:0 [120] R ==> qemu-system-aar:1781 [120]
qemu-system-aar-1781 [004] 83328.209649: kvm_irq_line:    Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1781 [004] 83328.209651: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1795 [005] 83328.209666: kvm_exit:      <CANT FIND FIELD exit_reason>IRQ: HSR_EC: 0x0000 (UNKNOWN), PC: 0xffffffffc0080f3298
qemu-system-aar-1795 [005] 83328.209672: kvm_entry:      PC: 0xffffffffc0080f3298
...
qemu-system-aar-1795 [005] 83328.209685: kvm_mmio:      mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

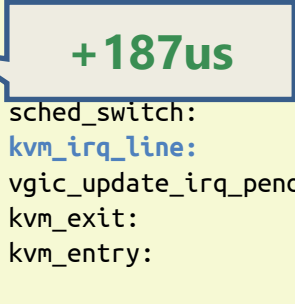
## vhost-net:



# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## vhost-net:

```
qemu-system-aar-1795 [005] 83328.209495: kvm_exit:      <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024 (DABT_LOW), PC: 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209496: kvm_guest_fault: ipa 0xa003000, hsr 0x93810046, hxfar 0xffffffffc009675a50, pc 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209498: kvm_mmio:      mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1795 [005] 83328.209512: sched_waking:    comm=vhost-1781 pid=1793 prio=120 target_cpu=001
    <idle>-0 [001] 83328.209522: sched_wakeup:        vhost-1781:1793 [120] success=1 CPU:001
qemu-system-aar-1795 [005] 83328.209525: kvm_entry:      PC: 0xffffffffc00872cd78
    <idle>-0 [001] 83328.209525: sched_switch:        swapper/1:0 [120] R ==> vhost-1781:1793 [120]
    vhost-1781-1793 [001] 83328.209533: kp_tun_get_user: (ffff800010b34550)
    vhost-1781-1793 [001] 83328.209568: sched_waking:    comm=qemu-system-aar pid=1781 prio=120 target_cpu=004
    vhost-1781-1793 [001] 83328.209572:              comm=vhost-1781 pid=1793 runtime=48480 [ns] vruntime=265274173669 [ns]
    vhost-1781-1793 [001] 83328.209574:              vhost-1781:1793 [120] S ==> swapper/1:0 [120]
    <idle>-0 [004] 83328.209577:              qemu-system-aar:1781 [120] success=1 CPU:004
    <idle>-0 [004] 83328.209584: sched_switch:        swapper/4:0 [120] R ==> qemu-system-aar:1781 [120]
    qemu-system-aar-1781 [004] 83328.209649: kvm_irq_line:  Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
    qemu-system-aar-1781 [004] 83328.209651: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
    qemu-system-aar-1795 [005] 83328.209666: kvm_exit:      <CANT FIND FIELD exit_reason>IRQ: HSR_EC: 0x0000 (UNKNOWN), PC: 0xffffffffc0080f3298
    qemu-system-aar-1795 [005] 83328.209672: kvm_entry:      PC: 0xffffffffc0080f3298
    ...
    qemu-system-aar-1795 [005] 83328.209685: kvm_mmio:      mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```





# Tracing Interruption Between Host and Guest: `trace-cmd -e kvm record`

## vhost-net:

```
qemu-system-aar-1795 [005] 83328.209495: kvm_exit:      <CANT FIND FIELD exit_reason>TRAP: HSR_EC: 0x0024 (DABT_LOW), PC: 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209496: kvm_guest_fault: ipa 0xa003000, hsr 0x93810046, hxfar 0xffffffffc009675a50, pc 0xffffffffc00872cd74
qemu-system-aar-1795 [005] 83328.209498: kvm_mmio:      mmio write len 4 gpa 0xa003a50 val 0x1
qemu-system-aar-1795 [005] 83328.209512: sched_waking:    comm=vhost-1781 pid=1793 prio=120 target_cpu=001
    <idle>-0 [001] 83328.209522: sched_wakeup:      vhost-1781:1793 [120] success=1 CPU:001
qemu-system-aar-1795 [005] 83328.209525: kvm_entry:      PC: 0xffffffffc00872cd78
    <idle>-0 [001] 83328.209525: sched_switch:      swapper/1:0 [120] R ==> vhost-1781:1793 [120]
vhost-1781-1793 [001] 83328.209533: kp_tun_get_user:    (ffff800010b34550)
vhost-1781-1793 [001] 83328.209568: sched_waking:    comm=qemu-system-aar pid=1781 prio=120 target_cpu=004
vhost-1781-1793 [001] 83328.209572:               comm=vhost-1781 pid=1793 runtime=48480 [ns] vruntime=265274173669 [ns]
vhost-1781-1793 [001] 83328.209574:               vhost-1781:1793 [120] S ==> swapper/1:0 [120]
    <idle>-0 [004] 83328.209577:               qemu-system-aar:1781 [120] success=1 CPU:004
    <idle>-0 [004] 83328.209584: sched_switch:      swapper/4:0 [120] R ==> qemu-system-aar:1781 [120]
qemu-system-aar-1781 [004] 83328.209649: kvm_irq_line:    Inject VGIC SPI interrupt (1), vcpu->idx: 0, num: 77, level: 1
qemu-system-aar-1781 [004] 83328.209651: vgic_update_irq_pending: VCPU: 0, IRQ 77, level: 1
qemu-system-aar-1795 [005] 83328.209666: kvm_exit:      <CANT FIND FIELD exit_reason>IRQ: HSR_EC: 0x0000 (UNKNOWN), PC: 0xffffffffc0080f3298
qemu-system-aar-1795 [005] 83328.209672: kvm_entry:      PC: 0xffffffffc0080f3298
...
qemu-system-aar-1795 [005] 83328.209685: kvm_mmio:      mmio unsatisfied-read len 4 gpa 0xa003a60 val 0x0
```

**+187us**

- Observed over 2 times time efficiency in vhost-net, as seen in the sequence trace
  - Highly consistent with bitrate measurements obtained using iperf3
- Essential reasons for this improvement would likely include:
  - Removal of qemu userspace (system call) from the critical path
  - Minimization of vCPU sleep & wakeup occurrences

# Agenda

- Introduction
- Debugging Tools for Virtualized Systems
- Practice: Analyzing vhost-net
- Summary



# Summary

- Confirmed the **value of vhost-net** by the behavior analysis using **specialized tools for virtualized environments**
  - Twice as fast as standard virtio-net
  - No significant side effects, especially in terms of CPU utilization
- **Various methodologies and tools are covered in this talk, they can be applied to a wide range of development scenarios for virtualized systems**
  - Analyzing performance improvements/regressions
  - Debugging/walking through virtualized systems

# Next Steps

- Contribute materials to AGL
  - Debugging tools
  - vhost-net support
- Explore additional tools
  - Port bcc scripts to aarch64 (e.g., kvmexit)
  - Develop custom tools designed for virtualized systems using bcc
  - Experiment with trace-cmd agent/listen
- Investigation on containerized and cloud environments

Thank You!

# Appendix

# Setting Up Tools on AGL: Installation Steps of Tools

```
cat << EOF >> conf/local.conf

# perf, trace-cmd
IMAGE_FEATURES += "debug-tweaks tools-debug tools-profile"

# bcc
IMAGE_INSTALL:append = " bcc"

# debuginfod
PACKAGECONFIG:append:pn-elfutils-native = "debuginfod libdebuginfod"
DISTRO_FEATURES:append = " debuginfod"

# tweak for debuginfod: we need to mask following recipe as it disrupts
# gdb's packageconfig and leads to a build error
BBMASK += "meta-qt5/recipes-devtools/gdb"

# ssh/sshfs
IMAGE_FEATURES:append = " ssh-server-dropbear"
IMAGE_INSTALL:append = " sshfs-fuse"

# iperf3
IMAGE_INSTALL:append = " iperf3"
EOF
```

# Setting Up Tools on AGL: Adding Kernel Configuration Fragments

- Enable kernel configs for tracing

```
$ cat << EOF >> conf/local.conf
IMAGE_INSTALL:append = " kernel-module-kheaders"
FILESEXTRAPATHS:prepend:pn-linux-renesas := "/path/to/put/kconfig:"
FILESEXTRAPATHS:prepend:pn-linux-yocto := "/path/to/put/kconfig:"
SRC_URI:append:pn-linux-renesas = " file://trace.cfg"
SRC_URI:append:pn-linux-yocto = " file://trace.cfg"
EOF
```

```
$ cat /path/to/put/kconfig/trace.cfg
trace.cfg
CONFIG_PERF_EVENTS=y
CONFIG_KPROBES=y
CONFIG_FTRACE=y
CONFIG_DYNAMIC_FTRACE=y
CONFIG_FUNCTION_TRACER=y
CONFIG_FUNCTION_GRAPH_TRACER=y
CONFIG_IRQSOFF_TRACER=y
CONFIG_PREEMPT_TRACER=y
CONFIG_SCHED_TRACER=y
CONFIG_FTRACE_SYSCALLS=y
CONFIG_TRACER_SNAPSHOT=y
CONFIG_PSTORE_FTRACE=y
CONFIG_BPF=y
CONFIG_BPF_SYSCALL=y
CONFIG_BPF_JIT=y
CONFIG_BPF_EVENTS=y
CONFIG_TASKSTATS=y
CONFIG_SCHEDSTATS=y
CONFIG_STACKTRACE=y
CONFIG_IKHEADERS=m
```

# Upgrading Tool Recipes

- Install the latest versions of `trace-cmd` and `bcc`
- Resolve issues related to building `bcc` against older kernel versions

```
# Add new layer
$ source agl-init-build-env
$ bitbake-layers create-layer ../meta-custom-tools
$ bitbake-layers add-layer ../meta-custom-tools
$ mkdir -p ../meta-custom-tools/recipes-devtools ../meta-custom-tools/recipes-kernel

# Clone latest recipes of the tools
$ git clone https://git.yoctoproject.org/poky /tmp/poky && git -C /tmp/poky checkout 34004afb65
$ git clone https://github.com/openembedded/meta-openembedded /tmp/meta-openembedded && \
  git -C /tmp/meta-openembedded checkout def4759e9
$ git clone https://github.com/kraj/meta-clang /tmp/meta-clang && git -C /tmp/meta-clang checkout 384dc8f

# Copy the updated recipes to the new layer
$ cp -r /tmp/meta-openembedded/meta-oe/recipes-kernel/trace-cmd ../meta-custom-tools/recipes-kernel
$ cp -r /tmp/meta-openembedded/meta-oe/recipes-kernel/libtracefs ../meta-custom-tools/recipes-kernel
$ cp -r /tmp/poky/meta/recipes-kernel/libtraceevent ../meta-custom-tools/recipes-kernel

$ cp -r /tmp/meta-clang/dynamic-layers/openembedded-layer/recipes-devtools/bcc ../meta-custom-tools/recipes-devtools
$ cp -r /tmp/meta-openembedded/meta-oe/recipes-kernel/libbpf ../meta-custom-tools/recipes-kernel
```

# Miscellaneous Setup

- Setting up vhost-net

On the host PC (before building):

```
$ cat << EOF >> conf/local.conf
PACKAGECONFIG:append:pn-qemu = " vhost"

IMAGE_INSTALL:append = " kernel-module-vhost-net"
FILESEXTRAPATHS:prepend:pn-linux-renesas := "/path/to/put/kconfig:"
SRC_URI:append:pn-linux-renesas = " file://vhost.cfg"
EOF

$ cat /path/to/put/kconfig/vhost.cfg
CONFIG_VHOST_NET=m
```

- Setting up debuginfod

On the host PC:

```
$ source /path/to/agl-init-build-env
$ oe-debuginfod
```

On target host system:

```
root@host:~# qemu-system-aarch64 \
    -enable-kvm \
    -cpu host \
    -netdev tap,helper="/usr/libexec/qemu-bridge-helper -
br=vmnet0",id=net0,vhost=on \
    [...]
```

On target (both host and guest):

```
root@host:~# export DEBUGINFOD_URLS=http://<ip-of-host-pc>:8002
```

```
root@guest:~# export DEBUGINFOD_URLS=http://<ip-of-host-pc>:8002
```